

Fintech and AI, read through the P&L: the margin sits with the licensed risk-holder run by software

Five fintech sub-sectors and the AI stack, one worked P&L each. The numbers first, then the claim, then what has to be true. Discussion document for JSW Ventures, 5 Oct 2026.

In the last thirty days the state reset two tolls and the market paid for one profit. UPI merchant payments carry a 0.4% fee above Rs 2,000, capped at Rs 300, from 15 October (NPCI). IRDAI proposed caps on what an insurance distributor may earn on 23 September, and PB Fintech lost 43% in four sessions. Moneyview listed 62% above issue on 1 October on a 14.44% cost of borrowings. Read through the ladder, these say one thing: revenue that is a toll on distribution is granted and withdrawn by rule and collected by licensed players at scale; revenue that is a yield or a premium survives if credit cost and claims are controlled. The investable company holds the risk on a licensed book and runs the operation with software instead of branches and agents. The AI argument is the same test from the other side: software earns the margin only where it holds the outcome.

What this rules out: standalone payments, consumer apps, broking, insurance and wealth distribution, unlicensed neobanks, per-seat software that makes a bank's process cheaper, horizontal AI tools.

The last thirty days	P&L line
15 Sep 2026 UPI merchant fee returns: 0.4% above Rs 2,000, capped at Rs 300, from 15 Oct	F1 revenue: a toll set by rule
23 Sep 2026 IRDAI consultation proposes caps on distributor commissions and insurer expenses; PB Fintech falls 43% in four sessions	F3 revenue: commission is a regulated toll
Sep 2026 NPCI presents the Unified Agent Protocol and an AI-agent registry for UPI: 'AI may recommend, but authentication and final settlement must follow deterministic auditable rules'	F1 and AI: when the buyer is software, the licensed layer keeps the margin
25 Sep 2026 JSW One Platforms files a DRHP for a Rs 3,054 cr IPO, Rs 500 cr for its lending arm	F2: embedded credit inside a workflow
1 Oct 2026 Moneyview lists 62% above issue after 98x subscription; cost of borrowings 14.44% in the DRHP	F2 and exits: the market pays for profit

The ladder I use on every segment, and where AI touches it

What is left after each cost, read top to bottom. CM1 to CM3 are my convention; companies define them differently, so ask which one a company means.

Line	Meaning	Quality looks like
Revenue	What the customer pays	Recur, grows with use, priced on value not cost; net of pass-through (interchange, model bill)
Gross margin	Less the direct cost of delivering it	Rises as volume grows; flat means a reseller
CM1	Less the variable cost of serving one unit	Positive on day one
CM2	Less the cost of keeping the customer: support, onboarding, collections, review	Falls per unit as software replaces people
CM3	Less the cost of winning the customer: sales, marketing, partner payouts	Each new customer pays back inside 18 months
EBITDA	Less fixed cost: technology, risk, compliance, management	Shows whether scale pays

Lender or insurer. The first test is risk-adjusted: yield less cost of funds less credit cost; premium less claims. Bajaj Finance borrows at 7.54%, Moneyview at 14.44%; the gap is the rating, and a model edge that a higher cost of funds eats earns nothing.

Pass-through businesses. Judge net revenue. A payment aggregator's gross revenue is mostly the bank's and the network's bill; an AI tooling company's gross revenue is mostly the model vendor's.

Where AI touches the ladder. Cost side (CM2): AI removes work in support, review, collections, and a cheaper model hands that saving to the buyer. Revenue side: AI helps the customer approve, convert, retain, and protects price. In my task map of 96 cells, 58 are coded cost-side and 40 revenue-side, 21 both; these are my codes, not data. The line that matters is who carries the outcome: a vendor paid per call, or an owner holding the loan.

Fintech: five sub-sectors, one worked unit each, one test

Does the sub-sector let a company that raises Rs 10 to 20 cr become a licensed risk-holder run by software, or sell a must-have to the ones that do?

Sub-sector	Worked unit (numbers)	What has to be true	Verdict (my reasoning)	Kill test	Exit evidence
F1 Payments	A Rs 10,000 UPI payment: fee Rs 40, acquirer keeps an assumed Rs 16; the average ticket (Rs 1,220) pays nothing	Net take rate 0.1 to 0.3% (assumption); Rs 1 cr of monthly net revenue needs Rs 333 to 1,000 cr of volume above the floor	Fails standalone; passes as a wedge inside vertical software for Rs 2,000 to 75,000 tickets	Payments above 40% of revenue by year three	M&A by listed acquirers; Pine Labs trades 22% below issue
F2 Lending	Rs 100 of loans: new lender earns -5.0% on assets in year one (-13.0% to 3.0%); Bajaj earns 5.2%	Cost of funds under 12%, opex under 4% of assets by year three, credit cost under 2.5%; then the same yield earns 4.5%	Passes in one asset class where repayment is controlled by the rail and the first Rs 100 cr of debt is contracted at entry	Opex to assets over 4% by year three; credit cost over 2.5%; no co-lending or DFI line in 12 months	IPO prices profit: Moneyview +62% on debut, Aye flat
F3 Insurance	A Rs 20,000 health policy: first-year commission Rs 3,500, half to marketing; year two Rs 1,200 at 80% persistency	Renewal share above 50% of three-year commission (today 38.2%); claims software at Rs 150 per claim against the insurer's Rs 1,000	Passes on claims and underwriting infrastructure and embedded covers; fails on distribution (IRDAI caps proposed 23 Sep)	Renewals under 80%; revenue tied to a commission rate the regulator is capping	PB Fintech -43% in four sessions; Acko and InsuranceDekho planning IPOs
F4 Wealth	Rs 5 lakh investor: trail Rs 3,250 a year, advice Rs 3,750, alternatives Rs 7,500; acquisition paid back in 33, 29, 14 months	Revenue at least 0.6% of assets, recurring, payback inside 18 months; execution is won (Groww Rs 735 cr Q1 PAT), F&O traders -19%	Passes only on manufacturing alternatives (SIF Rs 31,175 cr in 18 months) with captive distribution; fails on execution, MF distribution, advice	Revenue to assets under 0.6%; distribution rented from others	Groww bought Fisdrom (USD 150 mn); IPO
F5 Fraud, identity	One check: price Rs 10 (range Rs 2 to Rs 30), data fee Rs 5, gross margin 50% and rising with volume if the company owns the data	A rule creating the budget: banks pay 85% of coerced losses up to Rs 25,000 from 1 Jan 2027; DPDP 13 Nov 2026 and 13 May 2027; only 20.8% of regulated entities used AI in 2025	Passes for named modules with a regulatory deadline: fraud and mule detection, consent pipes, collections stacks, model-risk tooling	Top five customers over 50%; price per call falling over 20% a year	Perfios bought three companies in 2025; Setu bought Agya in 2026

Fintech P&Ls (1 of 2): payments and lending, with the numbers in the ladder

Payments is a regulated toll; lending is a capital-stack bet. The one line I would judge each on is in the amber box.

Payments and money movement (F1)

Revenue. MDR 0.4% on merchant payments above Rs 2,000, cap Rs 300 (binds at Rs 75,000), from 15 Oct 2026. On Rs 10,000: Rs 40. On the Rs 1,220 average ticket (NPCI Sept 2026: Rs 29.37 lakh cr over 24,068 mn): nothing. More than 95% of merchant transactions pay nothing (NPCI).

Who keeps it. Not published. Acquiring side assumed at 40% (30 to 50%): Rs 16 on the example payment. Net take rate after pass-through 0.1 to 0.3% (assumption).

Costs. Interchange, scheme and bank fees passed through; fraud and chargebacks 0.05% of value (assumption); risk, compliance, engineering.

Scale needed. Rs 1 cr of monthly net revenue needs Rs 333 to 1,000 cr of monthly volume above the floor. Paytm earns 9.0% PAT margin on Rs 2,448 cr of quarterly revenue, most of it lending distribution and devices; PhonePe lost Rs 2,792 cr in FY26.

Fundable? Only as a wedge inside vertical software for merchants with Rs 2,000 to 75,000 tickets, where the toll is largest in percentage terms.

Judge payments on net take rate and the share of revenue that is not volume. If net revenue grows faster than volume, the company sells something above the rail.

Lending and credit infrastructure (F2)

Ladder, Rs 100 of loans. New lender year one: yield 23%, cost of funds 13.5%, credit cost 4.5%, opex 10% of assets; return on assets -5.0% (-13.0% to 3.0%). Bajaj Finance: cost of funds 7.54% (FY26), ROA about 5.2%. Small-ticket listed lender: about 5.3%.

Capital stack. Rs 15 cr of equity supports a Rs 68 cr own book at 3.5x debt, or Rs 675 cr co-lent under the directions in force since 1 Jan 2026 (originator keeps 10%, default loss guarantee capped at 5%). Year-one loss on the own book: Rs 3.4 cr.

What has to be true. Cost of funds under 12% and opex under 4% of assets by year three, credit cost under 2.5% on a 12-month vintage. Then the same yield earns 4.5%. Moneyview borrows at 14.44% (DRHP); the first Rs 100 cr of debt, contracted at entry, decides mortality.

Sector. Unsecured retail GNPA 1.7% against 0.7% secured (RBI FSR, Mar 2026); gold loans doubled to Rs 5.14 lakh cr. Navi lost 3.6% of its book in FY26; slice, now a bank, made over Rs 50 cr in Q1 FY27.

Fundable? Yes, in one asset class where repayment is controlled by the rail (device locks, employer or institution deduction, UPI autopay on receivables), not by people. StrideOne is the in-house precedent.

Judge lending on credit cost by vintage against cost of funds. A model edge that a higher cost of funds eats earns nothing.

Fintech P&Ls (2 of 2): insurance, wealth, fraud and identity

Same ladder, different line decides each one.

Insurance (F3)

Distributor, Rs 20,000 health policy. First-year commission 17.5% (Rs 3,500), of which 50% goes to marketing and payouts; renewal 7.5% at 80% persistency gives Rs 1,200 in year two. Renewal share of three-year commission: 38%. All ranges are assumptions; the IRDAI paper of 23 Sep proposes to cap the first number.

Carrier. Premium less claims less expenses. Digit's Q1 FY27 profit fell 38% as claims rose; Acko turned profitable at Rs 43.6 cr in FY26. Individual health premiums grew 29.7% in Oct to Mar FY26 after GST went to zero.

Claims software. Rs 150 per claim (range Rs 50 to Rs 300) against the insurer's own Rs 1,000; price at 15% of the insurer's cost.

Fundable? Claims and underwriting infrastructure, embedded covers with retention economics. Not a distributor whose revenue is a commission rate under consultation.

Judge a distributor on renewal share of revenue; a carrier on loss ratio by cohort.

Wealth and capital markets (F4)

Rs 5 lakh investor. Trail 0.65% of assets (Rs 3,250 a year), advice 0.75% (Rs 3,750), alternatives 1.5% (Rs 7,500). Acquisition at Rs 9,000 pays back in 33, 29 and 14 months. Ranges are assumptions.

Market. Execution is won: Groww Rs 735 cr of Q1 FY27 profit and 86% above issue; Zerodha flat at a 51% margin. 87.7% of individual F&O traders lost money in FY26 and their number fell 19% (SEBI); STT up from 1 Apr 2026; MF Regulations 2026 halved brokerage caps. SIP Rs 32,297 cr a month; SIF AUM Rs 31,175 cr in 18 months.

What has to be true. Revenue at least 0.6% of assets, recurring, payback inside 18 months. No Indian firm found earning a profit on paid advice to the mass affluent.

Fundable? Manufacturing of alternatives (SIF, AIF, PMS) with captive distribution. Not broking, not MF distribution, not advice.

Judge wealth on revenue as a share of assets versus per trade; assets-based revenue survives a quiet market.

Fraud, compliance and identity (F5)

One check. Price Rs 10 (range Rs 2 to Rs 30), data fee Rs 5; gross margin 50%. A Rs 2 cr bank contract is 2.0 million checks a year.

The budget, by rule. Banks compensate 85% of coerced-transaction losses up to Rs 25,000 from 1 Jan 2027; DPD consent managers from 13 Nov 2026, full obligations 13 May 2027; mule-account procedure (SC Aug 2026, RBI draft Sep 2026); RBI model-risk draft (Jun 2026).

Buyers. Fewer than 200 regulated entities matter (assumption); 20.8% used or built AI in 2025 (RBI). Account Aggregator: 176 data providers, 1,076 data users, 538 mn consents fulfilled (Sahamati, Jul 2026). Credgenics: Rs 226 cr FY25 revenue.

Fundable? Yes, for modules a regulated entity must have on a date in the calendar, with net revenue retention above 1.2x and no customer above 20% of revenue. Resellers of others' data do not fit.

Judge fraud and identity on gross margin as checks grow; flat means a reseller.

AI: the layers, the markets, and the one test that survives 'is that not just efficiency again'

Compute and models are not fundable at this cheque; tooling and apps are, selectively. Beyond efficiency has one meaning that holds.

Layer or market	Revenue unit	Judge on	Fundable at Rs 10 to 20 cr?
Compute (AL1)	Price per GPU hour: IndiaAI portal Rs 65 to 100; Neysa raised USD 1,200 mn, half debt	Contracted utilisation against cost of capital	No
Foundation models (AL2)	Price per million tokens: USD 0.1 (small) to USD 2 (mid-tier); Sarvam raised USD 308 mn in 2026	Gross margin after inference as prices fall; top-customer share	No
Tooling: orchestration, data (AL3)	Platform fee plus a share of routed model spend; Databricks claims 35% routing saving (vendor, coding tasks)	Net revenue after pass-through model spend	Selective
Tooling: security, audit (AL4)	Subscription Rs 0.2 to 2 cr per regulated buyer (assumption); RBI model-risk draft creates the budget	Whether a rule or audit creates the budget	Yes
Consumer apps (AC)	2.2% of US households pay (a16z, Apr 2026); Emergent earns 8.5% of USD 120 mn in India	Paying share; inference per heavy user against price	Mostly no
Enterprise apps (AE, AF1 to AF4, AV1)	Per seat, per unit of work, or per outcome; see next slide	Price tied to the job, not to tokens; gross margin rising with scale	Yes, core fit
Physical (AP)	Revenue per deployed unit; payback 12 to 36 months (assumption)	Payback per unit; recurring share of revenue	Selective

The test. Efficiency: the same output at lower cost or time. The buyer keeps that saving, and a cheaper model hands more of it over every quarter. Per-outcome pricing is spreading (Intercom USD 0.99 per resolution, Agentforce USD 2 per conversation) and it is a contract; Genpact and EXL have sold gain-share for years without becoming software businesses. The one meaning of beyond efficiency that holds is that the software company becomes the operator and carries the outcome on its own account, visible as a loss ratio, a claims reserve or an SLA penalty in its own accounts, and the incumbent cannot follow because of a licence, data that holds through a cycle, or a structure it cannot cannibalise.

The Indian case. SquadStack moved its voice agents to a base fee plus a success fee on disbursed loans for DMI Finance (CIOL, 14 Sep 2026). If those loans go bad the loss sits with DMI. A success fee is a vendor's price; a loss is an owner's. The balance sheet shows which one a company is, and that is the diligence test.

Where the market is. Bessemer (Aug 2026), Stellaris (Aug 2026) and Elevation (Sep 2026) all back AI-native services priced per outcome; Hang Ten raised USD 85 mn in seed rounds for it. Direction agreed; the stopping point is too early. Bessemer's own warning: removing people does not by itself give software margins.

AI P&Ls: four worked units with the numbers

Voice and customer operations, finance agents, AI-native services, vertical AI. Human cost is the floor the price must sit under.

Customer operations and voice (AF2): one resolved call

Human. Agent at Rs 35,000 a month (Rs 25,000 to Rs 45,000), 115 calls a day: Rs 13.8 per call (Rs 7.6 to Rs 25.6). Assumption; the deck's v6 left this blank.

AI. 4,000 tokens at USD 0.68 per million (Ramp, Sep 2026) is Rs 0.24; telephony and speech Rs 2.5; total Rs 2.74 per call.

Price. Rs 8 per resolved call (Rs 5 to Rs 12): gross margin 66%, buyer's saving Rs 5.8 per call. Intercom charges USD 0.99 (Rs 86); India's cheap calls leave a thin price room.

Must be true. Resolution at or above the human benchmark; price per job, not per token; exception staff under 20% of revenue. Stellaris: 66% of 216 voice companies target an existing budget.

AI-native services (AF3): one finished job

Price. Per finished job against the services budget; Sequoia counts USD 6 of services spend per USD 1 of software (US). Hang Ten sells IT services with teams of 2 to 4 instead of about 30 (USD 85 mn raised).

Margin. 50% gross (40% to 60%), assumption. Bessemer: removing people does not give software margins.

Must be true. Human minutes per job fall every quarter; otherwise it is an agency with AI branding. The company that holds the outcome (prices the loss, not the task) is the one to fund.

Finance and back-office agents (AF1): one invoice

Price. Rs 20 per invoice (Rs 5 to Rs 50); Hyperbots sells an unlimited-access licence (Inc42). The pricing path that grows with the work is seat, then usage, then per document.

Costs. Inference Rs 0.35 per invoice; ERP integration Rs 0.3 cr per enterprise (0.1 to 0.5); first-year contract Rs 1.5 cr.

Margin. 98% gross per invoice before integration; integration is 20% of the first-year contract against a 20% threshold.

Must be true. A per-document price that holds as the agent replaces seats; integration cost that falls with each customer. If each customer is a project, it is a services firm.

Vertical AI (AV1): one site deployed

Price. Rs 8,000 per site a month (Rs 2,000 to Rs 20,000); deployment Rs 20,000 per site: 2.5 months of revenue to recover.

Must be true. Deployment repeats without a project (about one month of revenue or less), and the data gathered improves the product for the next site. Elevation (Sep 2026): Indian agents are 'judged on atoms', paired with on-the-ground fulfilment. Evidence in my set is thinnest here.

What I would fund at JSW: five licensed risk-holders run by software, plus two or three suppliers they cannot do without

Only the Mint facts about the fund, then arithmetic. The fund described needs low mortality and six outcomes at 5 to 8x; a 30x is not required.

Slot	What	Entry test (can fail)	Exit route
1, 2	Two licensed lenders: one secured retail class where repayment is controlled by the rail (EV fleets, education, LAP to GST filers); one embedded MSME lender that owns transaction data and settlement	First Rs 100 cr of debt contracted at entry; credit cost under 2.5% on a 12-month vintage; opex under 4% by year three	IPO on Rs 80 to 150 cr of PAT in year six or seven; Aye and Moneyview are the precedents
3	Claims and underwriting infrastructure for insurers, or an embedded cover with retention economics	Price per claim under one fifth of the insurer's own cost; renewals above 80%	Strategic or IPO; Acko and InsuranceDekho are filing
4	A manufacturer of alternatives (SIF, AIF, PMS) with captive distribution to the mass affluent and NRIs	Revenue above 0.6% of assets; acquisition paid back inside 18 months; distribution owned	Strategic (Groww and Fisdom precedent) or IPO
5	A licensed cross-border or outcome-holder (exporter collections with FX and credit; an AI company that holds the loss)	Take rate above 0.5%; a loss ratio line in its own accounts	Strategic
6 to 8	Suppliers to the above: fraud and mule detection, consent pipes, collections stacks, model-risk tooling	Net revenue retention above 1.2x; top five customers under 50%; a rule with a date behind the budget	M&A by Perfios-type acquirers and listed fintechs

Fund III target, midpoint

Rs 450 cr

Mint; not yet closed

Per name over life

Rs 30 cr
initial plus equal reserve

Needed back for 3x

Rs 1,350 cr

3.0x on every name if no zeros

The maths. Three zeros, six names at 1 to 2x and six at 5 to 8x return about 3.2x. A software-run lender listing on Rs 80 to 150 cr of PAT at 15 to 25x returns 5 to 8x on a diluted 8 to 10% stake. The 10x candidates are the suppliers with pricing power. Early-stage fintech funding was USD 367 mn in H1 2026 and seed USD 68.6 mn (Tracxn): a Rs 15 cr lead cheque sets price.

What would prove me wrong, by when. If by FY28 the best software-run lenders do not show opex under 4% of assets with credit cost within 50 basis points of the branch lenders, the operating-team claim fails. If RBI's AI rules come to require human sign-off on every credit decision, the economics revert to the branch model (the June 2026 draft asks for validation and kill switches, not sign-off). If the regulator grants distribution a durable margin again, the distribution rows change; in September it did the opposite.

Glossary and sources

Every term used, in plain English. Sources are listed in the workbook Inputs sheet with URL, date, claim type and confidence.

Term	Meaning
MDR	Merchant discount rate: the fee deducted from what a merchant receives on a card or UPI payment.
P2M, P2P	Person to merchant; person to person.
Take rate	Revenue as a share of the money that flows through.
Interchange	The slice of a card fee paid to the cardholder's bank; a US debit average of 0.73% in 2024 (Fed).
Yield	Interest and fees earned as a share of loans.
Cost of funds	Interest paid as a share of borrowings.
Credit cost	Loan losses and provisions as a share of loans.
Opex to assets	Operating cost as a share of the loan book.
ROA	Profit as a share of assets.
Vintage	Loans grouped by the month they were made, so losses can be compared like for like.
Co-lending	A bank funds most of a loan an originator sources; RBI directions from 1 Jan 2026 set 10% minimum retention and a 5% cap on the default loss guarantee.
Default loss guarantee	A promise by the originator to cover the first losses on a pool of loans.
GNPA	Gross non-performing assets: loans overdue more than 90 days as a share of all loans.
Persistency	The share of insurance policies renewed.
Loss ratio	Claims paid as a share of premium.
Trail	Recurring commission on assets a distributor keeps while the money stays invested.
SIF	Specialised investment fund: SEBI's category between mutual funds and PMS, Rs 10 lakh minimum.
STT	Securities transaction tax.
Account Aggregator	RBI-licensed consent system through which a customer shares bank and other data with a lender.
NRR	Net revenue retention: this year's revenue from last year's customers over last year's.
Token	A unit of text a model reads or writes; prices are quoted per million.
Inference	Running a trained model to produce an answer; the per-use cost.
Forward-deployed	An engineer who sits with the customer to connect the product to its systems.

Primary sources used. NPCI (MDR FAQ 15 Sep 2026; UPI statistics Oct 2026; app-wise data Aug 2026); RBI (Financial Stability Report 2026; co-lending directions; fraud liability directions Jun 2026; FREE-AI report Aug 2025; model-risk draft Jun 2026); SEBI (FY26 derivatives study); IRDAI (consultation 23 Sep 2026); AMFI; Sahamati; Federal Reserve (Reg II); company results (Paytm, Pine Labs, PB Fintech, Go Digit, Groww, Angel One, Bajaj Finance) and DRHPs (Moneyview); Lok Sabha replies; Budget 2026.

Secondary and claims. Business Standard, Inc42, Entrackr, TechCrunch, INDmoney, Outlook Business, Medianama, CIOL; Ramp AI Index (measured on Ramp customers); a16z State of Markets II (chart values); Sequoia, Bessemer, Stellaris, Elevation (investors' own posts, labelled as claims).

Assumptions. Every assumption carries a range on the Inputs sheet and is marked amber. Nothing on these pages is invented; where I could not source a figure it is blank or labelled.